

Healthcare (ML) project

January 20, 2023

1 Import Libraries

```
[1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
[2]: heart_data= pd.read_csv('cep1_dataset.csv')
```

```
[3]: heart_data.head()
```

```
[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[4]: heart_data.shape
```

```
[4]: (303, 14)
```

```
[5]: heart_data.info
```

```
[5]: <bound method DataFrame.info of
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	

1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6
..
298	57	0	0	140	241	0	1	123	1	0.2
299	45	1	3	110	264	0	1	132	0	1.2
300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]>

```
[6]: heart_data.isnull().sum()
```

```
[6]: age      0
sex        0
cp         0
trestbps   0
chol       0
fbs        0
restecg    0
thalach    0
exang      0
oldpeak    0
slope      0
ca         0
thal       0
target     0
dtype: int64
```

```
[7]: heart_data.describe()
```

```
[7]:
```

	age	sex	cp	trestbps	chol	fbs \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000

	restecg	thalach	exang	oldpeak	slope	ca \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

2 Checking distribution of the target variable

```
[8]: heart_data['target'].value_counts()
```

```
[8]: 1    165
      0    138
      Name: target, dtype: int64
```

```
[9]: X= heart_data.drop(columns='target', axis=1)
      Y= heart_data['target']
```

```
[10]: print(X)
```

```
      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0      63   1   3      145   233   1         0      150     0       2.3
```

1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6
..
298	57	0	0	140	241	0	1	123	1	0.2
299	45	1	3	110	264	0	1	132	0	1.2
300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal
0	0	0	1
1	0	0	2
2	2	0	2
3	2	0	2
4	2	0	2
..
298	1	0	3
299	1	0	3
300	1	2	3
301	1	1	3
302	1	1	2

[303 rows x 13 columns]

```
[11]: print(Y)
```

0	1
1	1
2	1
3	1
4	1
..	..
298	0
299	0
300	0
301	0
302	0

Name: target, Length: 303, dtype: int64

3 Splitting the data into train and test data

```
[12]: X_train, X_test, Y_train, Y_test= train_test_split(X, Y, test_size=0.  
    ↪2, stratify=Y, random_state=2)
```

```
[13]: print(X.shape)
```

```
(303, 13)
```

```
[14]: print(X_train.shape)
```

```
(242, 13)
```

```
[15]: print(X_test.shape)
```

```
(61, 13)
```

4 Model Training Logistic Regression

```
[16]: model= LogisticRegression()
```

```
[17]: model.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

```
[17]: LogisticRegression()
```

5 Model Evaluation

5.1 Accuracy Score

```
[18]: # Accuracy on train data  
X_train_prediction=model.predict(X_train)  
training_data_accuracy=accuracy_score(X_train_prediction, Y_train)
```

```
[19]: print('Accuracy on Training data:', training_data_accuracy)
```

Accuracy on Training data: 0.8512396694214877

```
[20]: # Accuracy on test data
X_test_prediction=model.predict(X_test)
test_data_accuracy=accuracy_score(X_test_prediction, Y_test)
```

```
[21]: print('Accuracy on Test data:', test_data_accuracy)
```

Accuracy on Test data: 0.819672131147541

5.2 Building a Predication System

```
[22]: input_data=(63, 1, 3, 145, 233, 1, 0, 150, 0, 2.3, 0, 0, 1)
```

```
[23]: # change the input data to a numpy array
input_data_as_numpy_array=np.asarray(input_data)
print(input_data_as_numpy_array)
```

```
[ 63.    1.    3.  145.  233.    1.    0.  150.    0.    2.3    0.    0.
   1. ]
```

```
[24]: # reshaping the numpy array as we are predicating for only on instance
input_data_resaped=input_data_as_numpy_array.reshape(1,-1)
```

```
[25]: print(input_data_resaped)
```

```
[[ 63.    1.    3.  145.  233.    1.    0.  150.    0.    2.3    0.    0.
   1. ]]
```

```
[26]: predication=model.predict(input_data_resaped)
print(predication)
```

```
[1]
```

/usr/local/lib/python3.7/site-packages/sklearn/base.py:451: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

"X does not have valid feature names, but"

```
[27]: if (predication[0]== 0):
        print('The Person does not have a Heart Disease')
    else:
        print('The Person has Heart Disease')
```

The Person has Heart Disease

```
[28]: input_data= (60,1,0,117,230,1,1,160,1,1.4,2,2,3)
input_data_as_numpy_array=np.asarray(input_data)
input_data_resaped=input_data_as_numpy_array.reshape(1,-1)
predication=model.predict(input_data_resaped)
print(predication)
if (predication[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')
```

[0]

The Person does not have a Heart Disease

/usr/local/lib/python3.7/site-packages/sklearn/base.py:451: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

"X does not have valid feature names, but"

[]: